

DocIntel User Manual

Document Intelligence Platform for Specialty Insurance *Powered by
Snowflake Cortex AI*

Table of Contents

1. [Getting Started](#)
 2. [Overview Dashboard](#)
 3. [Document Browser](#)
 4. [Document Search](#)
 5. [Review Queue](#)
 6. [Doc Intelligence Agent](#)
 7. [Portfolio Analytics](#)
 8. [Saved Items](#)
 9. [Document Detail View](#)
 10. [Administration](#)
 - [Pipeline Management](#)
 - [Ingestion Registry](#)
 - [Quality & Feedback](#)
 - [Cost Dashboard](#)
 - [Extraction Schema & Test Lab](#)
 - [Classification Labels](#)
 - [Lines of Business](#)
 - [Notifications](#)
 - [Configuration](#)
 - [Audit Log](#)
 - [User Management](#)
 11. [Help & Tips](#)
-

Getting Started

DocIntel is an AI-powered document intelligence platform built for specialty insurance operations. It automatically parses, classifies, and extracts structured data from insurance documents including policies, claims, endorsements, binders, certificates, and correspondence.

Access & Authentication

Access to DocIntel is controlled through Snowflake role-based entitlements:

- **DOC_INTEL_USER** — Standard user access (documents, search, review, analytics)
- **DOC_INTEL_ADMIN** — Full access including all administration features

Your visibility into documents is determined by your assigned role's permissions: - **Line of Business (LOB) access** — Controls which LOBs' documents you can see - **Document Type access** — Controls which document types are visible - **Page access** — Controls which application pages appear in your sidebar

Navigation

The left sidebar provides access to all features. It is organized into two sections:

Main Navigation: | Page | Description | | — | — | — | — | | Overview | Pipeline health and recent activity | | Document Browser | Browse and filter all indexed documents | | Document Search | AI-powered content search | | Review Queue | Approve or correct low-confidence documents | | Doc Intelligence | Conversational AI agent for document questions | | Portfolio Analytics | Charts, KPIs, and natural language data queries | | Saved Items | Your saved searches and bookmarked documents | | Help | In-app documentation and tips |

Administration (visible to admins): | Section | Pages | | — | — | — | — | | Operations | Pipeline Management, Ingestion Registry, Quality & Feedback, Cost Dashboard | | AI Config | Extraction Schema, Classification Labels, Lines of Business | | Management | Notifications, Configuration, Audit Log, User Management |

The Review Queue badge in the sidebar shows the count of documents pending review.

Overview Dashboard

Path: / (Home)

The Overview page provides an at-a-glance summary of your document intelligence platform.

What You See

- **Hero Banner** — Platform title with quick-access buttons to Search Documents and Ask a Question
- **Metrics Cards** — Four key indicators:
 - **Total Indexed** — Documents fully processed and available for search
 - **In Queue** — Documents currently being processed
 - **Failed** — Documents that encountered processing errors
 - **Total Seen** — All documents ever registered by the platform
- **Quick Questions** — Pre-built questions you can click to instantly query the AI agent:
 - “What exclusions apply to our D&O policies?”
 - “How many cyber claims are open this year?”
 - “Find coverage opinions from the last 90 days”
 - “Which policies expire in the next 60 days?”
- **Recently Indexed** — The last 8 documents processed, showing source folder, document type, and format

First-Run Onboarding (Admin Only)

If no documents have been indexed yet, admins see a setup banner with links to: - Configure the ingestion pipeline - Set up user access and roles

Document Browser

Path: /documents

The Document Browser is your primary interface for browsing, filtering, and managing all indexed documents.

Filtering Documents

Use the filter bar at the top to narrow results:

- **Document Type** — Dropdown to filter by classification (e.g., Policy, Claim, Endorsement)
- **Line of Business** — Dropdown to filter by LOB (e.g., D&O, Cyber, E&O, Marine)
- **Text Search** — Search across insured name, carrier, policy number, and content
- **Needs Review Only** — Checkbox to show only documents flagged for human review

The total document count updates as you apply filters.

Document Table

The table displays documents with these columns:

Column	Description
Checkbox	Select documents for batch actions
Type	Document classification (links to detail page)
Insured	Insured party name
LOB	Line of Business (color-coded badge)
Format	File format (PDF, TIFF, DOCX, etc.)
Status	Processing status
Confidence	AI confidence score (green ≥85%, amber ≥70%, red <70%)
Effective Date	Policy effective date
Expiry Date	Policy expiration date
Policy #	Policy number

Viewing Document Details

Click any row to open the **side detail panel** showing:

- Document type and warning indicators (low confidence, PII detected)
- Full metadata: Insured, Carrier, LOB, Coverage Form, Policy #, Claim #, dates, Status, Pages, Format, Confidence
- Document summary
- **View Document** button to open the original file preview

Correcting Attributes

From the side panel, scroll to “Correct an Attribute”:

1. Select the field to correct from the dropdown
2. Enter the correct value
3. Click **Submit Correction**

Correctable fields include: Insured Name, Carrier, Document Type, Policy Number, Claim Number, Line of Business, Coverage Form, Status, Effective Date, Expiration Date, Retroactive Date.

Batch Actions

1. Select documents using row checkboxes
2. A floating action bar appears showing the selection count
3. Available actions: **Archive Selected**
4. Click Apply to execute or Clear to deselect

Export

Click the **Export CSV** button in the header to download the current filtered view as a CSV file.

Pagination

Results display 50 per page. Use Previous/Next buttons to navigate.

Document Search

Path: /search

Document Search uses Snowflake Cortex Search to find documents by their content using natural language queries.

Performing a Search

1. Type your query in the search bar (e.g., “D&O claims with cyber exclusions”)
2. Click **Search** or press Enter
3. Results appear ranked by relevance with matching snippets highlighted

Using Filters

Click the **Filters** toggle to expand additional options:

- **Line of Business** — Filter results to specific LOBs
- **Document Type** — Filter by classification
- **Format** — Filter by file format (PDF, TIFF, etc.)
- **Status** — Filter by document status
- **Max Results** — Control result count (10, 20, 50, or 100)

Suggested Searches

When the search bar is empty, suggested searches are displayed: - D&O Claims - Cyber FNOL - Coverage Opinions - E&O Submissions - Reinsurance Certs - Expiring Policies

Click any suggestion to run that search instantly.

Working with Results

Each result shows:

- Document type with format badge
- Insured/Party name
- LOB (color badge)
- Status (color badge)
- Matching content snippet

Actions per result: - **Thumbs Up/Down** — Provide relevance feedback to improve future searches - **Bookmark** — Save the document for quick access later (appears in Saved Items)

Result Detail Panel

Click any result to open the detail panel showing: - Full metadata grid (Insured, Carrier, Coverage Form, dates, Limits, Policy #, Claim #, Confidence) - Document summary - Additional matching excerpts with page numbers - Links to the full document detail page and file preview

Saving Searches

1. Click the **Save** button (floppy disk icon) next to the search bar
2. Enter a name for the search
3. Choose visibility:
 - **Just Me** — Only you can see this saved search
 - **My Role** — All users with your role can see it
 - **Everyone** — All platform users can see it
4. Click **Save**

Saved searches preserve your query and all active filters for one-click re-use.

Scoped Search

When navigating to search from a specific document's Raw Text tab, search can be scoped to a single file. A banner shows the scoped filename with a "Clear scope" button to return to global search.

Review Queue

Path: `/review`

The Review Queue surfaces documents that need human review — typically those with low AI confidence scores or competing classification candidates.

Understanding the Queue

- **Scope Display** — Shows your LOB and document type access (you only see documents within your entitlements)
- **Progress Bar** — Visual indicator showing Pending / Approved / Skipped counts with completion percentage

Status Tabs

Switch between views: - **Pending** — Documents awaiting review (with count badge) - **Approved** — Documents you've already approved - **Skipped** — Documents you've skipped - **All** — Complete list regardless of status

Document Cards

Each card in the queue displays:

- Filename and source folder
- Document type badge
- LOB badge
- **Confidence bar** — Visual indicator with color coding:
 - Green ($\geq 75\%$) — High confidence
 - Amber ($\geq 50\%$) — Medium confidence
 - Red ($< 50\%$) — Low confidence
- **Review reason** — Why this document was flagged (e.g., "Low confidence: 62%")
- **Also considered** — Alternative document types the AI considered (helps you decide if the AI got it right)
- Key metadata: Insured, Policy #, Claim #

Taking Action

Three actions are available for each pending document:

1. **Approve** (green checkmark) — Confirms the AI's classification is correct. The document exits the queue.
2. **Skip** (gray button) — Removes the document from your queue without approving or correcting. Use when you're unsure or the document isn't in your area.
3. **Correct** (amber button with dropdown) — Opens a dropdown showing alternative document types. Select the correct type to:
 - Submit the correction
 - Auto-approve the document with the corrected classification
 - The "Open detail for manual correction..." link at the bottom opens the full detail page for complex corrections

Pagination

Results display 20 per page with Previous/Next navigation.

Doc Intelligence Agent

Path: /chat

The Doc Intelligence Agent is a conversational AI interface powered by Snowflake Cortex Agent. It can search documents, read content, query portfolio data, and answer questions that span both unstructured documents and structured analytics.

Starting a Conversation

Type your question in the input field and press Enter or click Send.

Example questions:

- "What exclusions apply to our D&O policies?"
- "How many cyber claims are open this year?"
- "Find coverage opinions from the last 90 days"
- "What's the total reserve exposure for open claims?"

Click any of the pre-built quick questions in the empty state to get started immediately.

How It Works

The agent has access to multiple tools: - **Document Search** — Finds relevant documents by content - **Portfolio Analytics** — Queries structured data via Cortex Analyst

While the agent is working, you'll see animated status messages indicating which tool it's using (e.g., "Searching documents..." or "Analyzing data...").

Reading Responses

- Responses render as formatted Markdown including tables and lists
- **Tool indicators** show which tool was used (e.g., "via Search" or "via Analytics")
- **Charts** may be rendered inline when the agent returns data visualizations

Citations

Document references appear as clickable blue chips: - `[filename.pdf]` — Click to open the document preview - `[filename.pdf, p.5]` — Click to open the document at a specific page

Providing Feedback

Each assistant response has thumbs up/down buttons. Use these to rate response quality — this feedback helps improve the system over time.

Managing Conversations

- Click **New conversation** to clear the chat history and start fresh
- The conversation is maintained per session only

Deep Linking

You can pre-populate a question using the URL parameter `?q=`. This is how the Overview page's quick questions work.

Portfolio Analytics

Path: `/analytics`

Portfolio Analytics provides visual dashboards and a natural language query interface for analyzing your document portfolio.

KPI Cards

Four summary metrics at the top: - **Total Documents** — Count of all indexed documents - **Document Types** — Number of active taxonomy types - **Lines of Business** — Specialty lines being tracked - **Documents by LOB** — Active coverage areas with documents

Charts

Four visualization panels:

1. **Documents by Type** — Horizontal bar chart showing the top 10 document types by volume
2. **Line of Business** — Donut/pie chart breaking down documents by LOB with legend
3. **Documents by Line of Business** — Vertical bar chart with color per LOB
4. **Document Status** — Vertical bar chart with color coding:
 - Cyan = Open
 - Gray = Closed
 - Green = Bound
 - Amber = Quoted
 - Red = Declined
 - Purple = Reserved

Natural Language Queries

The analytics page includes a powerful natural language query interface powered by Cortex Analyst:

1. Type your question in the input field (e.g., “How many D&O claims are open by month this year?”)
2. Click **Analyze**
3. Results display:
 - **Prose answer** — Natural language summary with a copy button
 - **Generated SQL** — Toggle “Show generated SQL” to see the underlying query (with copy button)
 - **Results table** — First 10 rows of data displayed; truncated with total row count if more

This uses a Semantic View to translate your question into accurate SQL against the document portfolio.

Saved Items

Path: `/saved`

Saved Items stores your saved searches and bookmarked documents for quick re-access.

Saved Searches Tab

Displays all searches you've saved from the Document Search page.

Each entry shows: - Search name - Query text (truncated) - Created date

Actions: - **Run** (green play icon) — Re-executes the search with all original filters - **Copy Link** — Copies a shareable URL to your clipboard - **Delete** (red trash) — Permanently removes the saved search

Bookmarks Tab

Displays all documents you've bookmarked from Search or the Document Browser.

Each entry shows: - Document type - Insured name - Personal note (if added) - Created date - File path link (opens document preview)

Actions: - **Delete** (trash icon) — Removes the bookmark

Document Detail View

Path: `/documents/detail?id=<file_id>`

The Document Detail page provides a comprehensive view of a single document with full metadata, history, and raw text.

Header

- File name
- Badges: Document Type, LOB, Confidence Score, Status
- **Preview** button — Opens the original document in a viewer
- **Archive** button (admin only) — Hides the document from search results
- **Purge** button (admin only) — Permanently deletes the document and source file (requires double confirmation)

Navigation

Context-aware back navigation: - “← Back to Review Queue” (if you came from Review) - “← Back to Search” (if you came from Search) - “← Back to Documents” (if you came from the Browser) - Generic back button otherwise

Attributes Tab

Displays all extracted fields organized by category:

Editable fields (click the pencil icon to correct): - Insured Name - Carrier - Document Type - Policy Number - Claim Number - Line of Business - Coverage Form - Status - Effective Date - Expiration Date - Retroactive Date

Fields with pending corrections show a “Pending” badge.

Read-only fields: - Broker - Per Occurrence Limit - Aggregate Limit - Retention/Deductible - Summary - Exclusions Noted - Action Required

History Tab

Shows the processing version timeline: - Version number - Processing status (Complete, Failed, Processing) - Start and end timestamps - Number of attempts - Error details with the failing stage (if applicable)

Related Tab

Displays documents related to the current one (sharing the same insured or policy number). Each related document links to its own detail page.

Raw Text Tab

Shows the full extracted text content:

- **Page count** displayed at the top
- **Highlight Values** toggle — When enabled, color-codes extracted values found in the text:
 - Blue = Identifiers (policy #, claim #)
 - Green = Parties (insured, carrier, broker)
 - Orange = Dates (effective, expiration, retroactive)
 - Purple = Monetary (limits, retention)
 - Gray = Classification (document type, LOB)
 - Includes a color legend

- **Search within document** link — Opens Document Search scoped to this file
 - **Preview Original** button — Opens the source file viewer
-

Administration

Administration pages are available to users with the DOC_INTEL_ADMIN role.

Pipeline Management

Path: `/admin/pipeline`

Manages the end-to-end document processing pipeline with six tabs.

Overview Tab

- **Health Metrics** — Cards showing Complete, Pending, Processing, and Failed document counts
- **Stale Queue Warning** — Amber banner if queue age exceeds 120 minutes
- **Task Management** — View all Snowflake tasks with their state and schedule
 - **Suspend** — Pause a running task
 - **Resume** — Restart a suspended task

Upload Tab

Upload documents directly into the processing pipeline:

1. Drag-and-drop files onto the upload zone, or click to browse
2. Supported formats: PDF, TIFF, DOCX, PPTX, JPEG, PNG, HTML, TXT (max 50MB per file)
3. Click **Upload and Register** to push files to the ingestion stage

Stages Tab

Manage ingestion stages (Snowflake stages configured as document sources):

- View all registered stages with Active/Paused status
- **Toggle** individual stages between Active and Paused
- **Remove** a stage from the registry

- **Add New Stage** — Manual entry with stage name, description, and path prefix
- **Discover Available Stages** — Automatically scan the database for unregistered stages and add them

Run Pipeline Tab

Manually trigger individual pipeline steps: 1. **Register Files** — Scan stages for new files and add to the ingestion registry 2. **Parse Documents** — Run AI document parsing (OCR/text extraction) 3. **Classify** — Run AI classification on parsed documents 4. **Extract** — Run AI attribute extraction on classified documents

Each step runs independently. Use this for testing or recovering from partial failures.

Processing Logs Tab

View recent processing activity: - Columns: File, Status, Folder, Type, Insured, Error message - **Filter by status**: Pending, Processing, Complete, Failed, Abandoned - **Refresh** to see latest entries

Reprocess Tab

Re-queue documents for reprocessing:

1. **Search/filter** by filename, status, LOB, or document type
2. **Select** individual documents via checkboxes, or use “Select All”
3. **Reprocess Selected** — Queue chosen documents for reprocessing
4. **Reprocess All Filtered** — Bulk reprocess everything matching current filters
5. **Manual entry** — Enter a specific file path to reprocess

Ingestion Registry

Path: `/admin/registry`

The central record of all documents the platform has seen, regardless of processing status.

Table columns: | Column | Description | | — — | — — — — | | File Path | Full stage path (clickable to preview) | | Format | File type with colored badge | | Folder | Source folder/stage | | Status | Pending, Processing, Complete, Failed, or Abandoned | | Version | Processing version number | | Attempts |

How many times processing was tried | | First Seen | When the file was first discovered | | Last Error | Most recent error message (if any) |

Features: - **Filter** by status and filename search - **Reprocess** button per document (prompts confirmation for non-failed documents) - **Pagination** — 50 per page with Previous/Next - **Refresh** button

Quality & Feedback

Path: /admin/quality

Monitor extraction quality and manage user-submitted corrections.

KPI Cards

- **Average Confidence** — Mean AI confidence score across all documents
- **Low Confidence Count** — Documents below the confidence threshold
- **Low Conf %** — Percentage of documents with low confidence
- **Document Types** — Count of active document types

Charts

- **Corrections by Field** — Bar chart showing total corrections vs. approved corrections per field name
- **Average Confidence by LOB** — Bar chart with confidence scores and document counts per Line of Business

Pending Corrections Table

Displays corrections submitted by users that need admin review: - File name (clickable to preview) - Field corrected - Original value - Corrected value - Submitted by (username) - Date

Actions: - **Approve** — Accept the correction and update the document - **Reject** — Dismiss the correction

Recent User Feedback Table

Shows thumbs up/down feedback submitted from Search and Chat: - Feedback type (positive/negative icon) - File name - Query text that triggered the feedback - User - Date

Cost Dashboard

Path: `/admin/cost`

Track Snowflake credit usage attributed to the DocIntel platform.

KPI Cards

- **Total Credits (30 days)** — Credits consumed in the last 30 days
- **Daily Average** — Average daily credit consumption
- **Projected Monthly** — Extrapolated monthly cost at current rate

Charts

- **Daily Credit Usage** — Line chart spanning 1 year showing daily credit consumption
- **Credits by Stage** — Horizontal bar chart showing which pipeline stages consume the most credits
- **Credits by Event Type** — Vertical bar chart breaking down costs by operation type

Cost Optimization Tips

Best practice recommendations for reducing credit usage.

Extraction Schema & Test Lab

Path: `/admin/extraction-schema`

Configure what data the AI extracts from documents and test the extraction pipeline.

Schema Configuration Tab

The extraction schema uses a three-tier hierarchy:

1. **Common** — Attributes extracted from ALL documents regardless of type
2. **Category** — Attributes extracted only when a document matches a specific classification category
3. **LOB-Specific** — Attributes extracted when a document matches both a category AND a Line of Business

Navigating the schema: - Use the sidebar to browse tiers and groups - Each group shows its “Fires When” match rule (for Category and LOB tiers) - The attributes table shows: Name, Description, Data Type

Managing attributes: - **Add Attribute** — Create a new extraction field: - Select tier (Common, Category, or LOB) - Choose or create a group - Set match rules (which document types/LOBs trigger extraction) - Define field name, description, and type (VARCHAR, DATE, NUMBER, VARIANT) - **Delete** — Remove an attribute from the schema - **Test** (flask icon) — Test a single attribute against a pinned file to verify extraction works

Test Lab Tab

The Test Lab lets you run the complete pipeline on sample files without affecting production data.

File Selection: 1. Select a Snowflake stage from the dropdown 2. Browse available files (search by name, filter by LOB/type) 3. **Pin** individual files or **Pin All Filtered** to add to your test set 4. Pinned files persist in your browser (localStorage)

Running Tests: 1. Click **Run Pipeline** to execute: Parse → Classify → Detect LOB → Resolve Schema → Extract 2. Results display per file: - Classification result - LOB detection - All extracted fields with tier labels - Check icon (✓) for successfully extracted values - X icon (✗) for empty/missing values - Fields are sorted with found values first

Upload: Use the **Upload File** button to add test documents directly to a stage.

Results are NOT saved to the production pipeline — the Test Lab is purely for validation.

Classification Labels

Path: `/admin/classification-labels`

Manage the document type taxonomy used by the AI classifier.

Viewing Labels: - Labels are organized by category (e.g., Underwriting, Policy, Claims, Certificates) - Each label shows its active/disabled status

Managing Labels: - **Add Label** — Enter a label name and select/create a category - **Create Category** — Add a new grouping category - **Rename Category** — Edit a category name - **Delete Category** — Removes the category and all its labels - **Toggle Active/Disabled** — Disabled labels are

excluded from AI classification - **Delete Label** — Remove an individual label

Lines of Business

Path: `/admin/lobs`

Manage the Lines of Business (LOBs) that documents are assigned to.

Table columns: Code, Name, Description, Active status

Actions: - **Add LOB** — Enter a unique code, display name, and description - **Edit** — Modify the name or description inline - **Toggle Active/Disabled** — Disabled LOBs aren't used in classification - **Delete** — Remove a LOB (with confirmation dialog)

Notifications

Path: `/admin/notifications`

Configure automated notifications for pipeline events.

Viewing Rules: - Table shows channel type, endpoint, subscribed event types, and active status

Adding a Rule: 1. Choose channel: **Webhook** (HTTP POST to a URL) or **Email** (via Snowflake integration) 2. Enter endpoint/address 3. Select event types to subscribe to: - INGEST, PARSE, CLASSIFY, EXTRACT - ALERT, ARCHIVE, PURGE - OVERRIDE, FEEDBACK, DOWNLOAD 4. Set threshold (0 = notify on every event; higher numbers batch)

Managing Rules: - **Toggle** rule active/inactive - **Delete** a rule permanently

Configuration

Path: `/admin/config`

Manage platform-wide settings organized into logical groups.

Pipeline Settings

Setting	Description
<code>pipeline_enabled</code>	Master on/off for automated processing

Setting	Description
<code>batch_size</code>	Number of documents processed per task execution
<code>max_retries</code>	Maximum retry attempts for failed documents
<code>reprocess_after_days</code>	Auto-reprocess documents after N days
<code>max_processing_attempts</code>	Hard cap on total processing attempts
<code>pipeline_target_lag_minutes</code>	Target freshness for dynamic tables

AI Models

Setting	Description
<code>cortex_model</code>	Model for document classification
<code>extraction_model</code>	Model for attribute extraction
<code>agent_model</code>	Model powering the chat agent
<code>search_embedding_model</code>	Model for search embeddings (locked)

Model settings show a dropdown with available Cortex models. Some settings are locked (indicated by a lock icon) and cannot be changed through the UI.

Quality & Alerts

Setting	Description
<code>confidence_threshold</code>	Minimum confidence before flagging for review
<code>classify_confidence_threshold</code>	Minimum classification confidence
<code>alert_email</code>	Email address for system alerts
<code>admin_email</code>	Administrator contact email
<code>cost_alert_multiplier</code>	Cost spike threshold (e.g., 2.0 = alert at 2× normal)
<code>error_rate_threshold_pct</code>	Error rate % that triggers alerts

Each setting has an individual **Save** button. Changes take effect immediately.

Audit Log

Path: `/admin/audit`

Track all user and system actions for compliance and troubleshooting.

Filters: - **Time Range** — 7, 14, 30, or 90 days - **Event Type** — Filter by specific action type - **User** — Filter by username

Summary Cards: - Top event types with occurrence count and unique user count

Events Table: - Timestamp - User who performed the action - Event type - Affected file path

Export: Click **Export CSV** to download the filtered event list for external analysis.

User Management

Path: `/admin/users`

Manage who can access the platform and what they can see.

Users Tab

View and manage user entitlements:

- **Add User** — Select from Snowflake role members not yet registered, or enter manually; assign roles
- **Edit** — Change a user's role assignments
- **Deactivate** — Suspend a user's access without deleting their record
- **Reactivate** — Restore a deactivated user
- **Sync from Roles** — Automatically create entitlements for all Snowflake role members not yet configured

Roles Tab

Define application roles that control access:

Permission	Options
App Pages	Dashboard, Search, Review Queue, Ask a Question, Analytics, Document Browser, Saved Searches, Help
LOB Access	All (wildcard) or select specific LOB codes
Document Types	All (wildcard) or select from grouped categories
Downloads	Allow or deny document downloads

Actions: - **Create Role** — Define a new role with name, description, and permissions - **Edit** — Modify an existing role's permissions - **Delete** — Remove a role (reassign users first)

Lines of Business Tab

Manage LOBs available for role assignment: - Add, edit, deactivate, or set a default LOB - Separate from the AI Config LOBs page — this controls access entitlements

Help & Tips

Path: `/help`

Quick Tips

- **Search is natural language** — Type questions like “find D&O policies with cyber exclusions” rather than keyword-only queries
- **Follow up in chat** — The Doc Intelligence Agent maintains conversation context. Ask follow-up questions to refine answers.
- **Click citation chips** — Blue chips like `[policy.pdf, p.3]` open the source document at the referenced page
- **Batch actions in Browser** — Use checkboxes to select multiple documents for bulk operations
- **Bookmark from Search** — Click the bookmark icon on any search result to save it for later
- **Correct as you go** — Use the pencil icon in the Document Detail view to submit corrections. These improve the AI over time.

- **Keyboard shortcuts** — Press Enter to submit search queries and chat messages

Understanding Confidence Scores

Confidence scores represent the AI's certainty in its classification: - **≥85% (Green)** — High confidence, likely correct - **70–84% (Amber)** — Moderate confidence, may need review - **<70% (Red)** — Low confidence, flagged for human review

Documents below the configured threshold automatically appear in the Review Queue.

Document Processing Pipeline

Documents flow through these stages: 1. **Ingest** — File discovered on a stage and registered 2. **Parse** — AI extracts text content (OCR for images/ scans) 3. **Classify** — AI determines document type 4. **Extract** — AI pulls structured attributes based on the schema

If any stage fails, the document is marked as Failed and can be reprocessed from Pipeline Management.

Glossary

Term	Definition
LOB	Line of Business — the insurance specialty area (e.g., D&O, Cyber, E&O, Marine)
Document Type	The classification category assigned by AI (e.g., Policy, Claim FNOL, Endorsement, Binder)
Confidence	The AI's self-reported certainty (0–100%) in its classification
Stage	A Snowflake stage (cloud storage) where source documents are stored
Ingestion Registry	The master record of all files discovered by the platform
Extraction Schema	The configured fields that the AI attempts to extract from each document
Cortex Agent	Snowflake's AI agent framework powering the chat interface

Term	Definition
Cortex Analyst	Snowflake's natural language to SQL engine powering analytics queries
Cortex Search	Snowflake's AI-powered vector search for document content
RCR	Restricted Caller's Rights — SPCS security model where the app runs with limited permissions
Semantic View	A business-level data model that Cortex Analyst uses to translate questions to SQL

DocIntel v1.0 — Document Intelligence for Specialty Insurance Powered by Snowflake Cortex AI